

## Code

My code is based on the matrix multiplication code provided for the assignment. I modified it by adding a total of 6 *algorithms*, or variations on the order in which the matrix elements are addressed, and a 7th algorithm of my own design. I will describe why order matters and the results of my own algorithm in the **Results** section.

In addition to the provided code, I also wrote a Python script that ran the program for each algorithm using increasing matrix sizes to determine each algorithm's scalability. The script timed a multiplication of two matrices for each matrix size. Further, it did this 4 times and averaged the time to compute the product. It then increased the size of the matrix by a factor of 2 for the next set of multiplications. The scaling of the factors increased until a multiplication took longer than 10 minutes.

## Results

The fastest ordering of the operations was **jki** by a slim margin over **kji**. The order of **ikj** roughly tied with the ordering **kij** for the distinction of the worst ordering of operations. See Fig. 1 and the raw data at the end of the paper for details of each algorithm.




Fig. 1: Timing of multiplication algorithms

When considering the number of floating operations per second each algorithm was able to manage, the most- and least-efficient maintain their titles. Both the **jki** and **kji** orderings show initial high rates of calculations, followed immediately by sharp increases in efficiency and an eventual plateauing as the matrix sizes increase. The two least efficient algorithms start at about half the rate of the leaders and decrease quickly as the sizes increase. See Fig. 2.

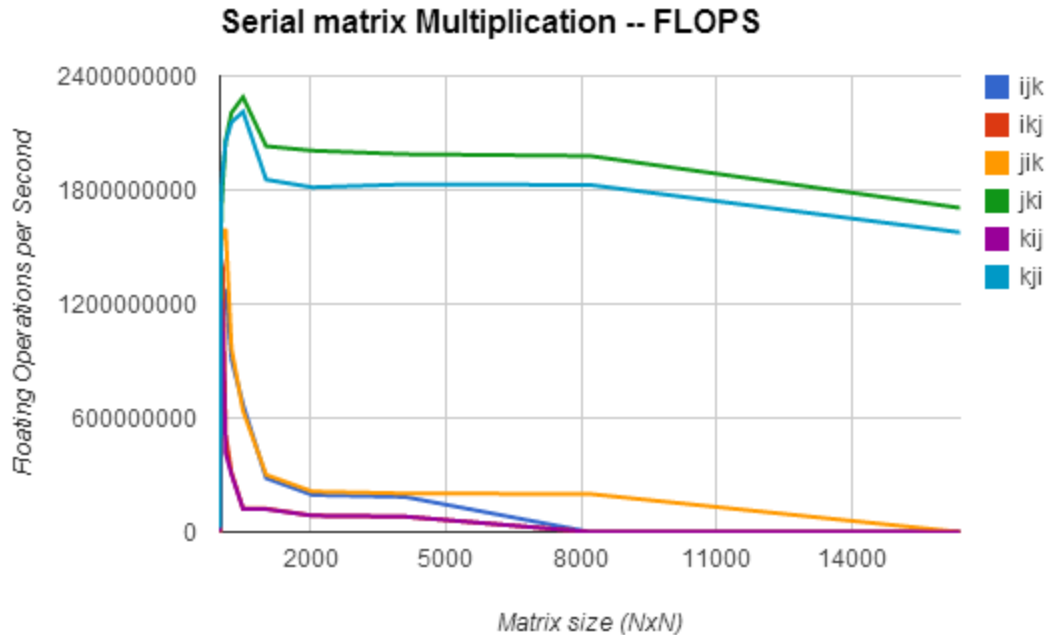


Fig. 2: Multiplication rate in floating operations per second

## Conclusion

When we analyze the algorithms that performed the best we find that the key is in how the data is accessed. When matrix elements are accessed in the same order in which they are laid out, specifically **column-major** order, we are able to leverage memory locality to our advantage. We get the benefit of temporal locality when we repeatedly access the same memory locations in consecutive operations. We are also able to benefit from spatial locality since data is read from main memory into cache not at the byte level but as blocks. This means that not only the element needed for the current operation will be cached, but it is very likely that the adjacent elements that we get as well will be needed for calculations in the immediate future.

For my own algorithm, I attempted to come up with a way to better take advantage of the above localities. After much whiteboard work and experimentation, I came up with an algorithm that worked in times almost identical to the **kji** ordering. Upon comparison I realised that the reason for this was that they *were* the same algorithm. So short of using a faster processor with more cache, I think the best algorithms determined through experimentation are near-ideal for a serial implementation.

The Python script was helpful for performing the analysis of the algorithms. It allowed me to run about 7 hours of timing runs on a CSIL machine overnight while it was unburdened. Without it I would have been less accurate in my record keeping and certainly less rested.

## Data

	2	4	8	16	32	64	128	256	512	1024	2048	4096	8192	16384
i		0.000		0.000			0.003	0.036						
j		0007	0.000	0072	0.000	0.000	3102	9412	0.399	7.583	88.15	743.84		
k	0	5	0015	5	053	4085	5	5	823	983	5881	158825		
i		0.000	0.000	0.000	0.000		0.008		2.194			1700.1		
k		0002	0012	0057	0457	0.000	1757	0.106	3957	17.54	199.3	413677		
j	0	5	5	5	5	3755	5	7945	5	06335	54981	5		
j				0.000	0.000									
i		0.000	0.000	0062	0427	0.000	0.002	0.034	0.420	7.139	80.69	678.09	5498.8	
k	0	0005	0015	5	5	335	628	741	3455	64125	8704	673075	62167	
j				0.000	0.000	0.000			0.117					
k	0.000	0.000	0.000	0067	0387	2862	0.002	0.015	2682	1.057	8.560	69.123	555.55	5156.6
i	0005	001	002	5	5	5	0365	2055	5	6595	155	02775	6624	57473
k					0.000	0.000		0.110		17.86	199.7	1695.7		
i		0.000	0.000	0.000	0537	4977	0.010	4412	2.198	62697	60307	413732		
j	0	0005	0015	007	5	5	1265	5	5845	5	5	5		
k	0.000	0.000		0.000	0.000	0.000	0.002							5577.8
j	0007	0012	0.000	0062	0362	2757	0417	0.015	0.121	1.158	9.471	75.125	601.91	260432
i	5	5	002	5	5	5	5	5615	378	06975	27825	9675	240225	5